
Virtual Memory

Chapter 18

S. Dandamudi

Outline

- | | |
|---|--|
| <ul style="list-style-type: none">• Introduction• Virtual memory concepts<ul style="list-style-type: none">* Page replacement policies* Write policy* Page size tradeoff* Page mapping• Page table organization<ul style="list-style-type: none">* Page table entries• Translation lookaside buffer | <ul style="list-style-type: none">• Page table placement<ul style="list-style-type: none">* Searching hierarchical page tables• Inverted page table organization• Segmentation• Example implementations<ul style="list-style-type: none">* Pentium* PowerPC* MIPS |
|---|--|

Introduction

- Virtual memory deals with the main memory size limitations
 - * Provides an illusion of having more memory than the system's RAM
 - * Virtual memory separates logical memory from physical memory
 - » Logical memory: A process's view of memory
 - » Physical memory: The processor's view of memory
 - * Before virtual memory
 - » Overlaying was used
 - It is a programmer controlled technique

Introduction (cont'd)


- Virtual memory
 - * Automates the overlay management process
 - » Big relief to programmers
- Virtual memory also provides
 - * Relocation
 - » Each program can have its own virtual address space
 - » Run-time details do not have any impact on code generation
 - * Protection
 - » Programs are isolated from each other
 - A benefit of working in their own address spaces
 - » Protection can be easily implemented

Introduction (cont'd)

- Principles involved are similar to the cache memory systems
 - * Details are quite different
 - » Due to different objectives
 - * Concept of locality is key
 - » Exploits both types of locality
 - Temporal
 - Spatial
 - * Implementation is different
 - » Due to different lower-level memory (disk)
 - Several orders of magnitude slower

Virtual Memory Concepts

- Implements a mapping function
 - * Between virtual address space and physical address space
- Examples
 - * PowerPC
 - » 48-bit virtual address
 - » 32-bit physical address
 - * Pentium
 - » Both are 32-bit addresses
 - But uses segmentation



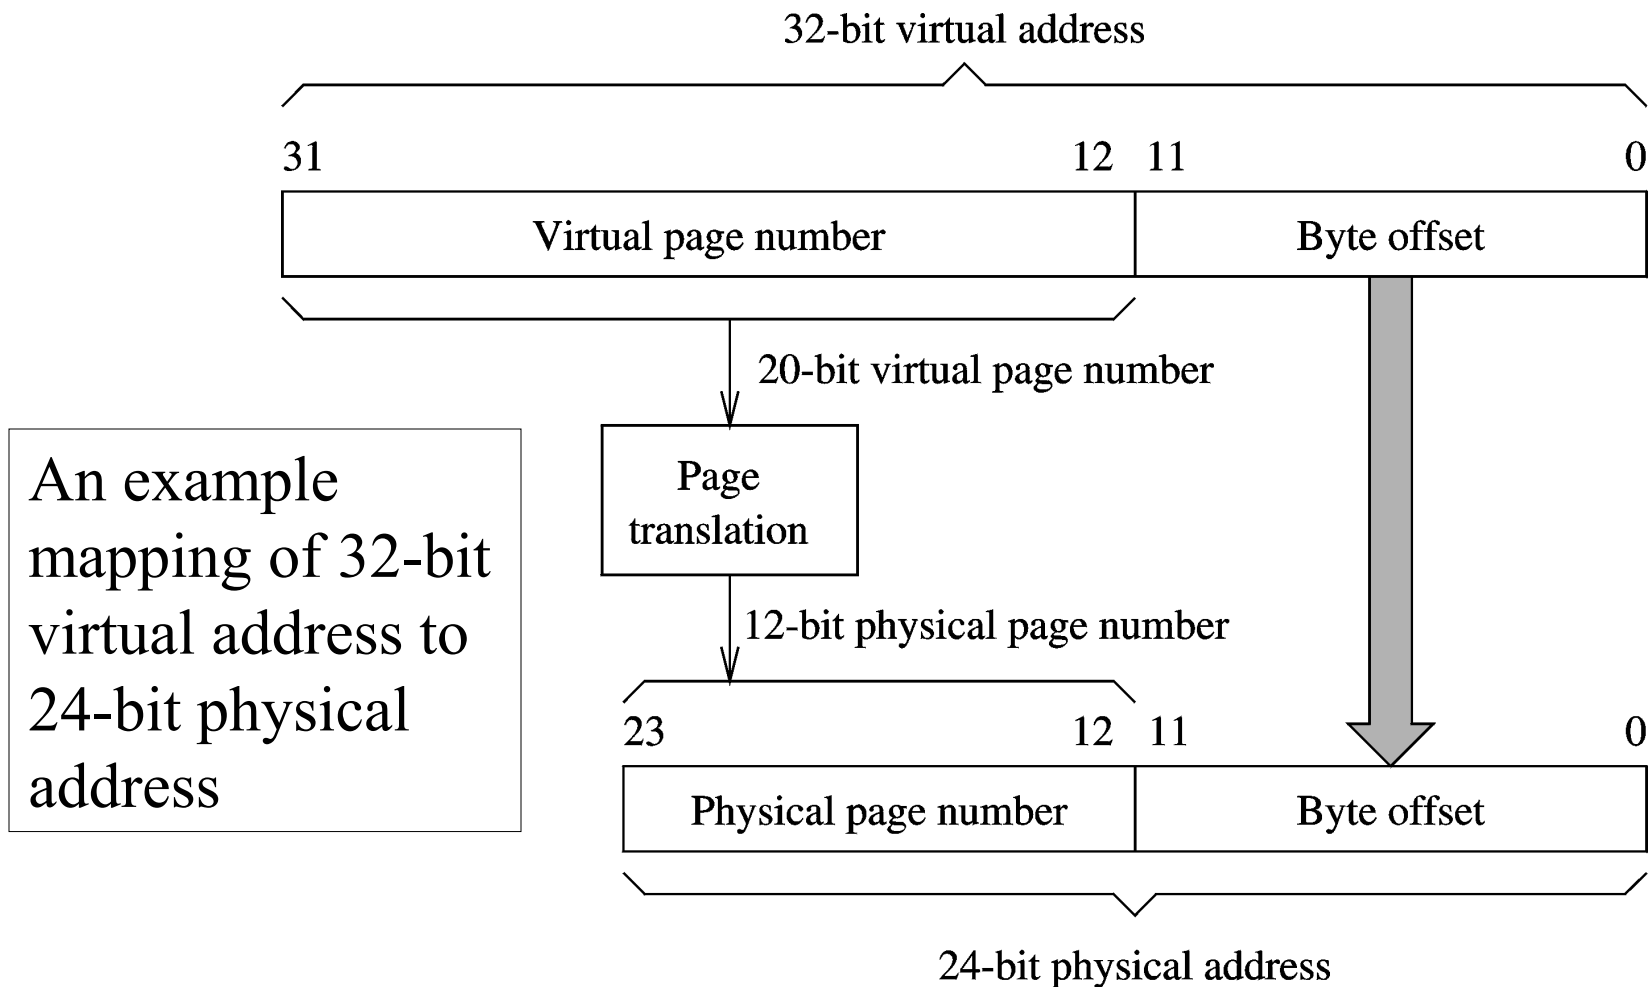
Virtual Memory Concepts (cont'd)

- Virtual address space is divided into fixed-size chunks
 - * These chunks are called *virtual pages*
 - * Virtual address is divided into
 - » Virtual page number
 - » Byte offset into a virtual page
 - * Physical memory is also divided into similar-size chunks
 - » These chunks are referred to as *physical pages*
 - » Physical address is divided into
 - Physical page number
 - Byte offset within a page

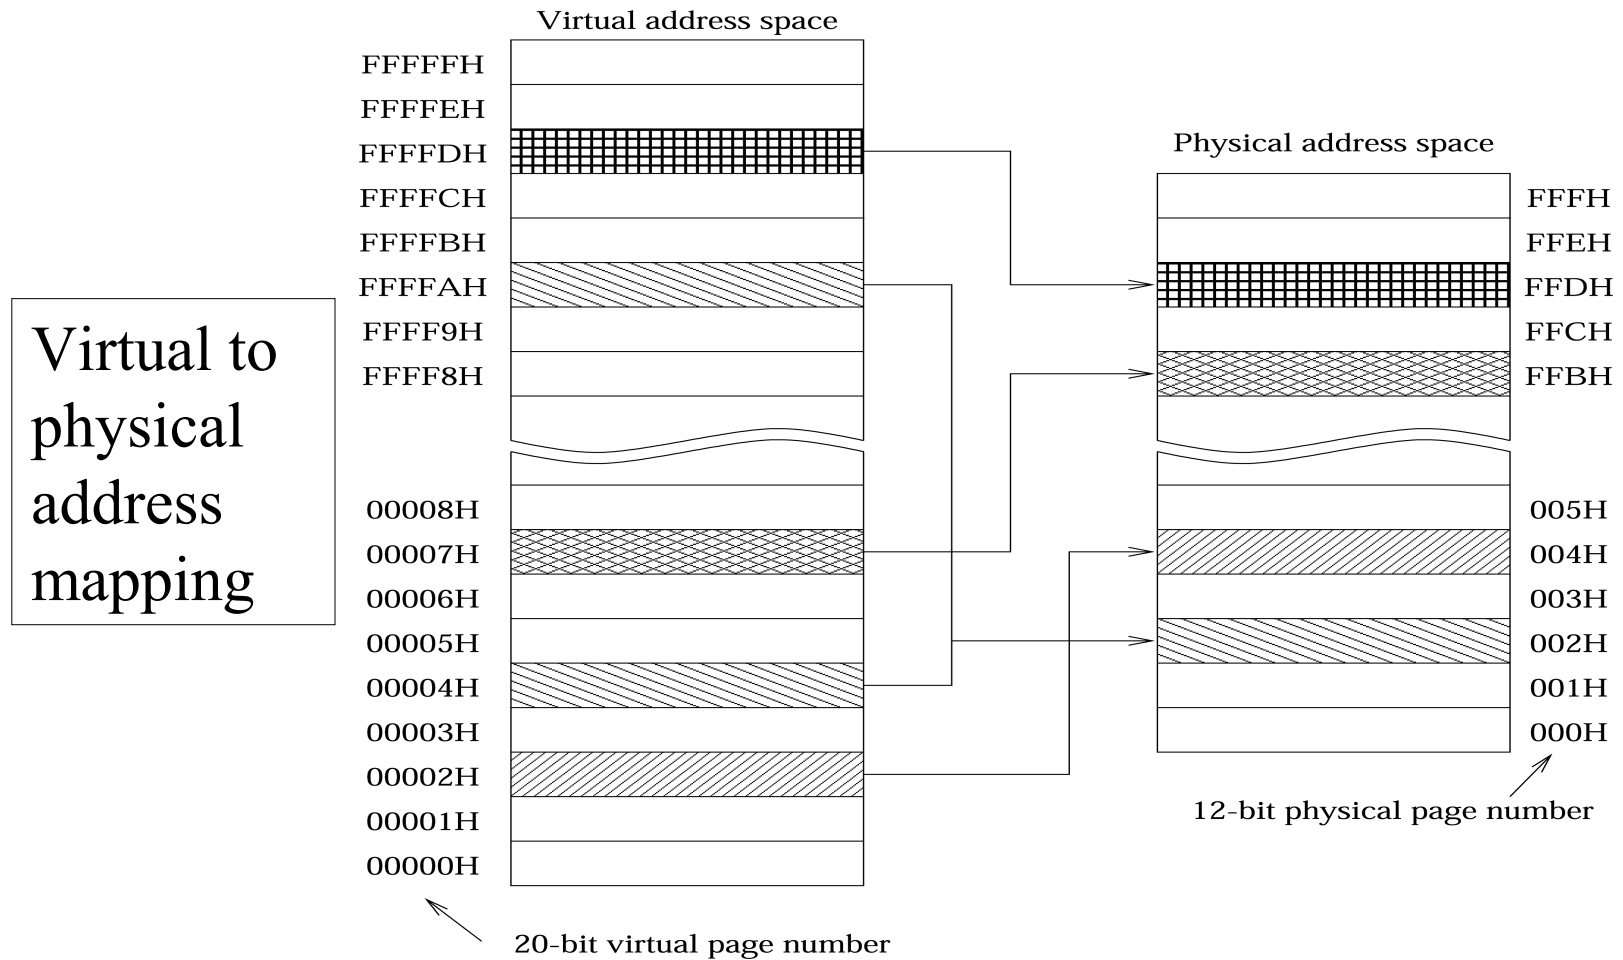
Virtual Memory Concepts (cont'd)

- Page size is similar to cache line size
- Typical page size
 - » 4 KB
- Example
 - * 32-bit virtual address to 24-bit physical address
 - * If page size is 4 KB
 - » Page offset: 12 bits
 - » Virtual page number: 20 bits
 - » Physical page number: 12 bits
 - * Virtual memory maps 2^{20} virtual pages to 2^{12} physical pages

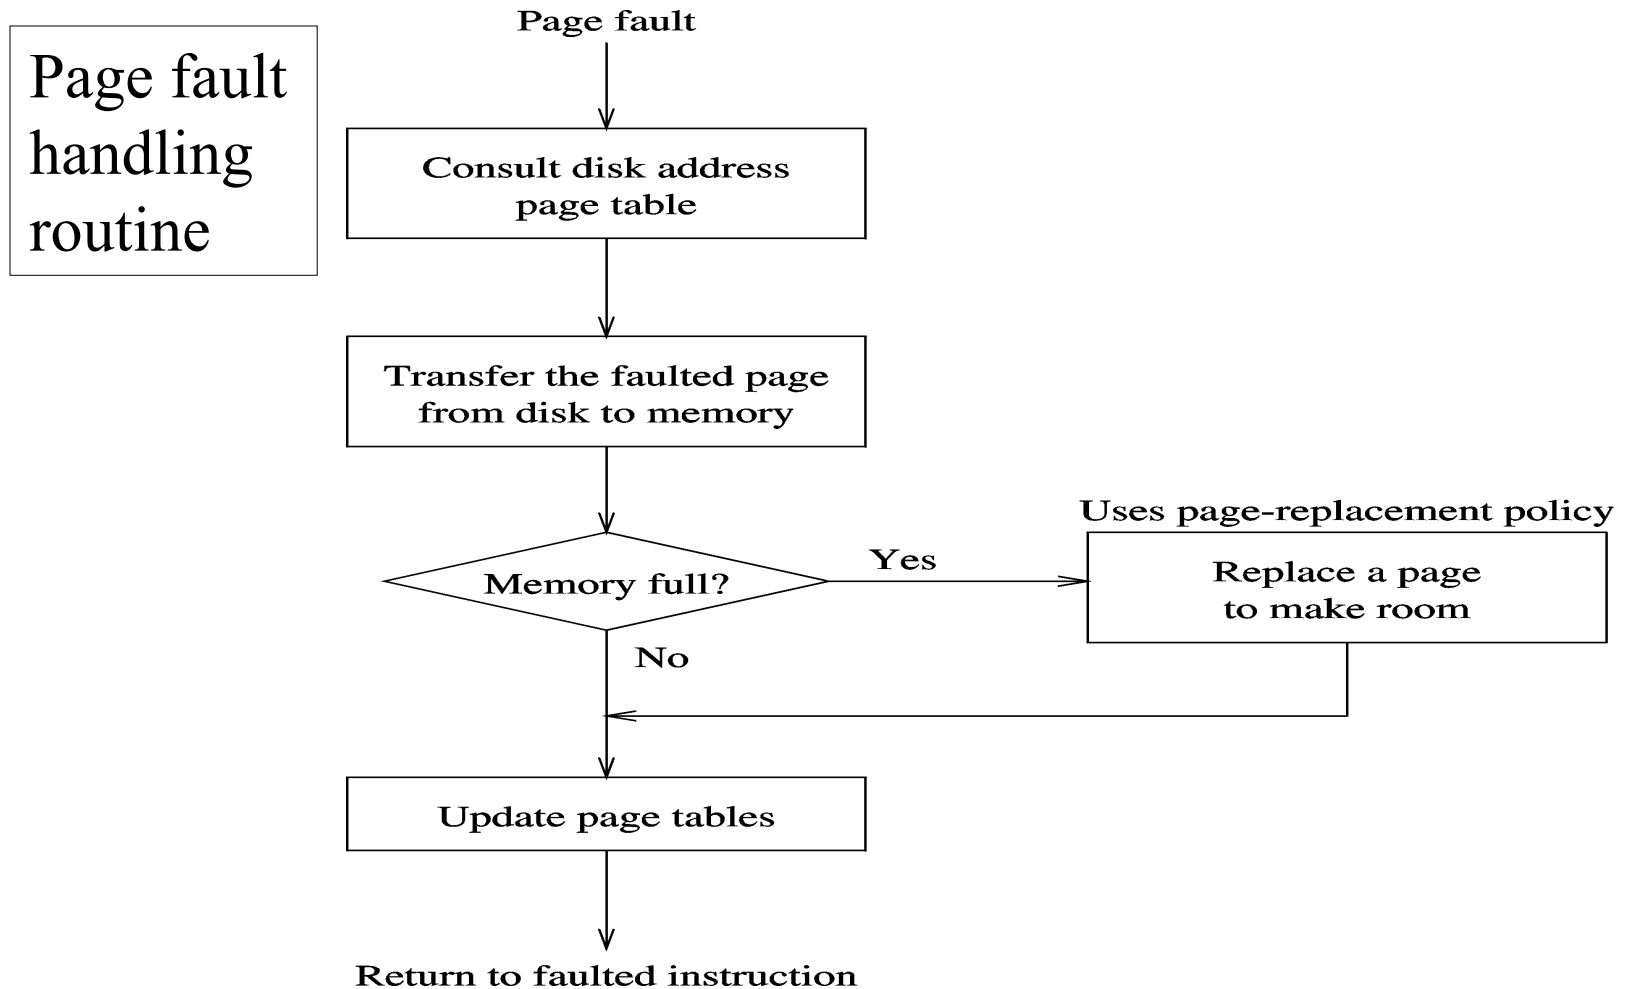
Virtual Memory Concepts (cont'd)



Virtual Memory Concepts (cont'd)

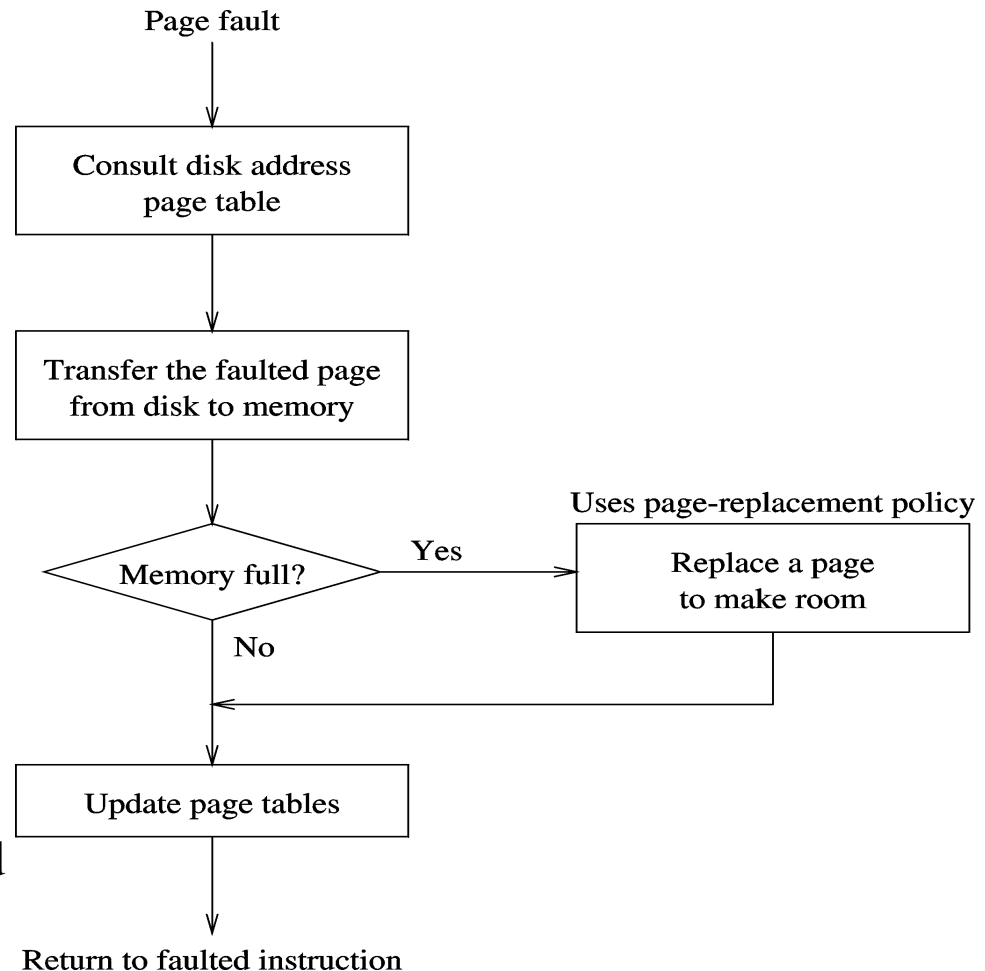


Virtual Memory Concepts (cont'd)



Virtual Memory Concepts (cont'd)

- A virtual page can be
 - * In main memory
 - * On disk
- Page fault occurs if the page is not in memory
 - * Like a cache miss
- OS takes control and transfers the page
 - * Demand paging
 - » Pages are transferred on demand



Virtual Memory Concepts (cont'd)

- Page Replacement Policies
 - * Similar to cache replacement policies
 - * Implemented in software
 - » As opposed to cache's hardware implementation
 - * Can use elaborate policies
 - » Due to slow lower-level memory (disk)
 - * Several policies
 - » FIFO
 - » Second chance
 - » NFU
 - » LRU (popular)
 - Pseudo-LRU implementation approximates LRU

Virtual Memory Concepts (cont'd)

- Write policies
 - * For cache systems, we used
 - » Write-through
 - Not good for VM due to disk writes
 - » Write-back
- Page size tradeoffs
 - * Factors favoring small page sizes
 - » Internal fragmentation
 - » Better match with working set
 - * Factors favoring large page sizes
 - » Smaller page sizes
 - » Disk access time

Pentium, PowerPC: 4 KB
MIPS: 7 page sizes between 4 KB to 16 MB

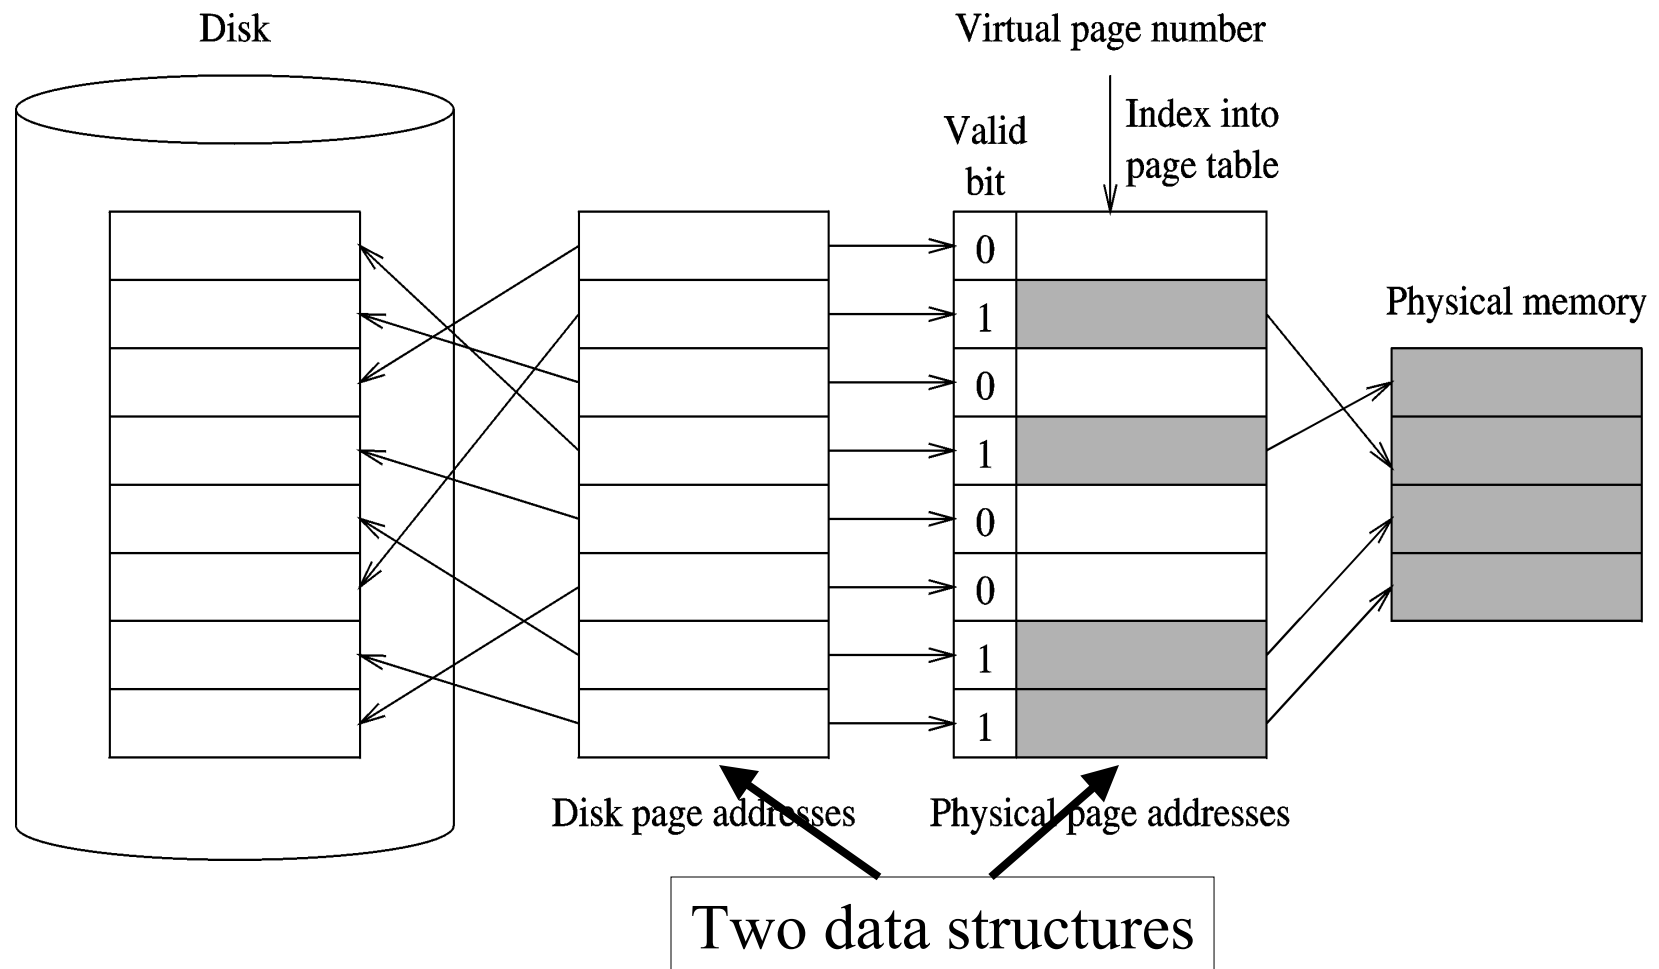
Virtual Memory Concepts (cont'd)

- Page mapping
 - * Miss penalty is high
 - » Should minimize miss rate
 - * Can use fully associative mapping
 - » Could not use this for cache systems due to hardware complexity
 - * Uses a translation table
 - » Called *page table*
 - * Several page table organizations are possible

Page Table Organization (cont'd)

- Simple page table organization
 - * Each entry in a page table consists of
 - » A virtual page number (VPN)
 - » Corresponding physical page number (PPN)
 - * Unacceptable overhead
- Improvement
 - * Use virtual page number as index into the page table
- Typical page table is implemented using two data structures

Page Table Organization (cont'd)



Page Table Organization (cont'd)

- Page table entry
 - * Physical page number
 - » Gives location of the page in memory if it is in memory
 - * Disk page address
 - » Specifies location of the page on the disk
 - * Valid bit
 - » Indicates whether the page is in memory
 - As in cache memory
 - * Dirty bit
 - » Indicates whether the page has been modified
 - As in cache memory

Page Table Organization (cont'd)

- * Reference bit
 - » Used to implement pseudo-LRU
 - OS periodically clears this bit
 - Accessing the page turns it on
- * Owner information
 - » Needed to implement proper access control
- * Protection bits
 - » Indicates type of privilege
 - Read-only, execute, read/write
 - » Example: PowerPC uses three protection bits
 - Controls various types of access to user- and supervisor-mode access requests

Translation Lookaside Buffer

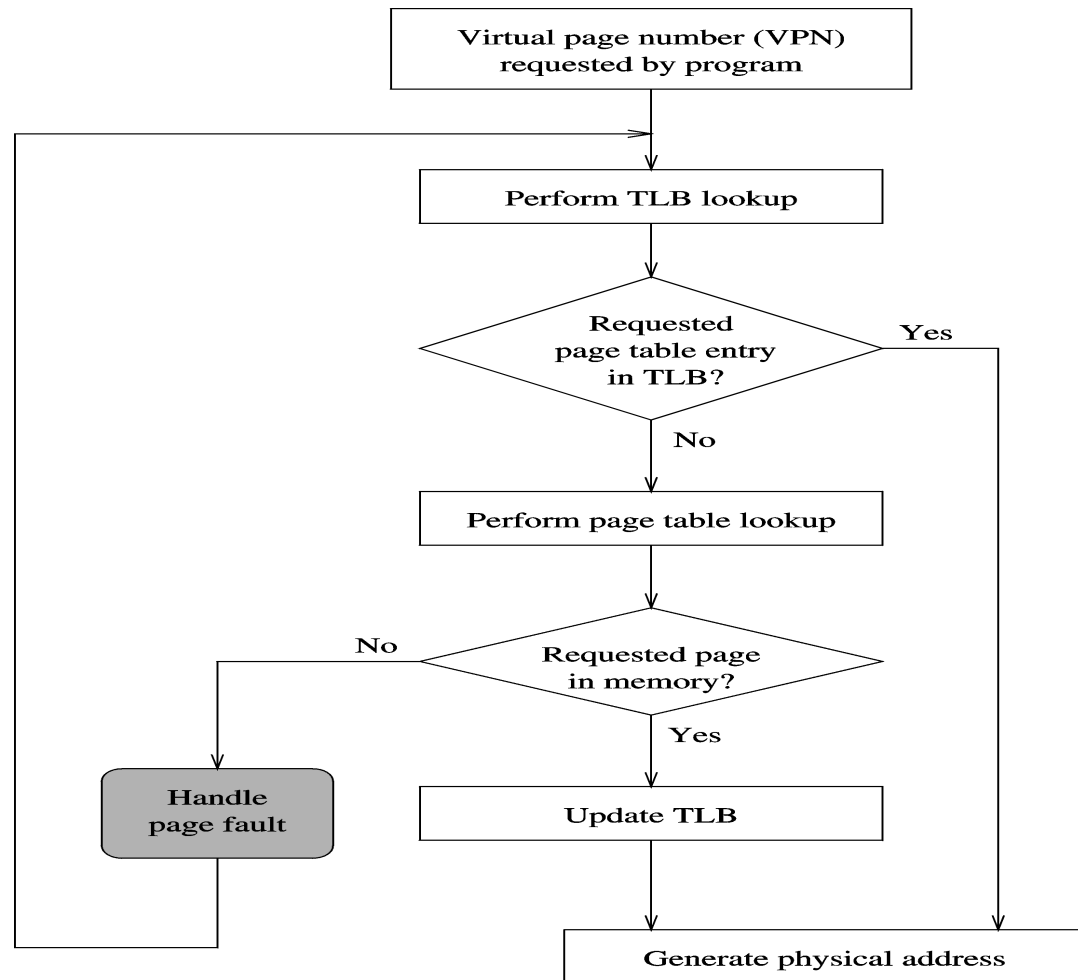
- For large virtual address spaces
 - * Translation table must be stored in virtual address space
 - » Every address translation requires two memory accesses:
 - To get physical page number for the virtual page number
 - To get the program's memory location
- To reduce this overhead, most recently used PTEs are kept in a cache
 - * This is called the *translation lookaside buffer* (TLB)
 - » Small in size
 - » 32 – 256 entries (typical)

Translation Lookaside Buffer (cont'd)

- Each TLB entry consists of
 - * A virtual page number
 - * Corresponding physical page number
 - * Control bits
 - » Valid bit
 - » Reference bit
 - » . . .
- Most systems keep separate TLBs for data and instructions
 - * Examples: Pentium and PowerPC

Translation Lookaside Buffer (cont'd)

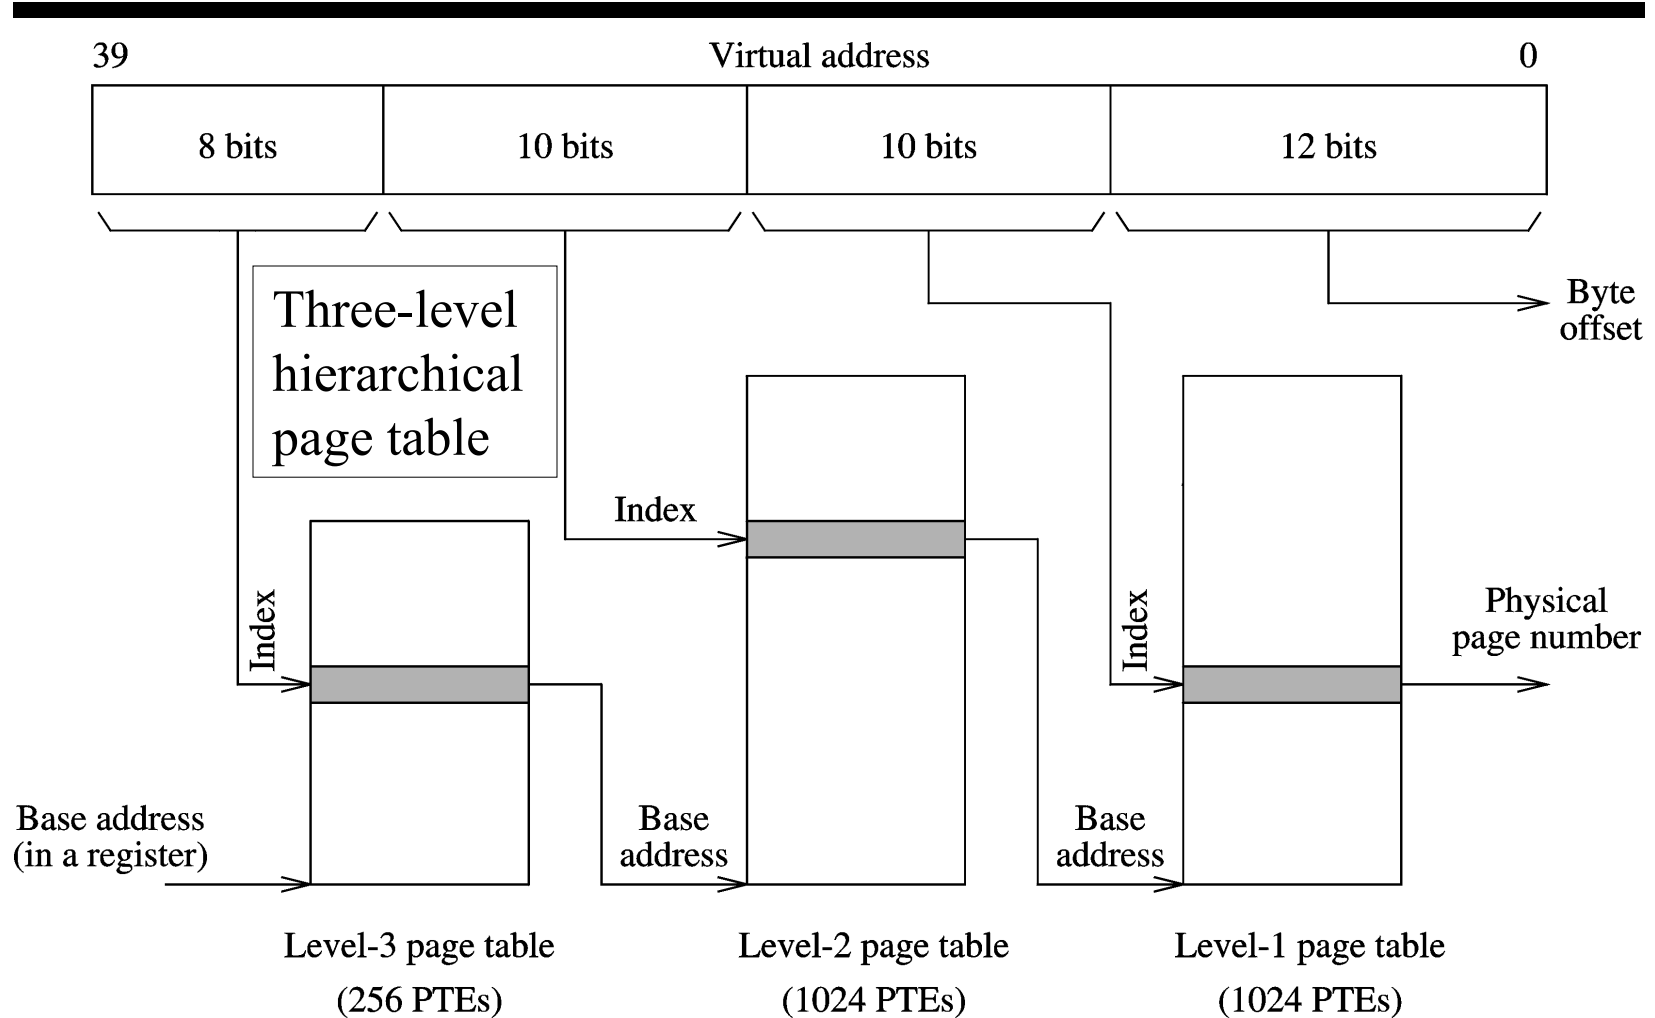
Translation
using a TLB



Page Table Placement

- Large virtual address spaces need large translation tables
 - * Placed in virtual address space
 - * Since the entire page table is not needed, we can use a hierarchical design
 - » Partition the page table into pages (like the user space)
 - » Bring in the pages that are needed
 - » We can use a second level page table to point to these first-level tables pages
 - » We can recursively apply this procedure until we get a small page table

Page Table Placement (cont'd)



Page Table Placement (cont'd)

- Hierarchical page tables are also called *forward-mapped page tables*
 - * Translation proceeds from virtual page number to physical page number
 - » In contrast to inverted page table
- Examples
 - * Pentium
 - » 2-level hierarchy
 - » Details later
 - * Alpha
 - » 4-level hierarchy

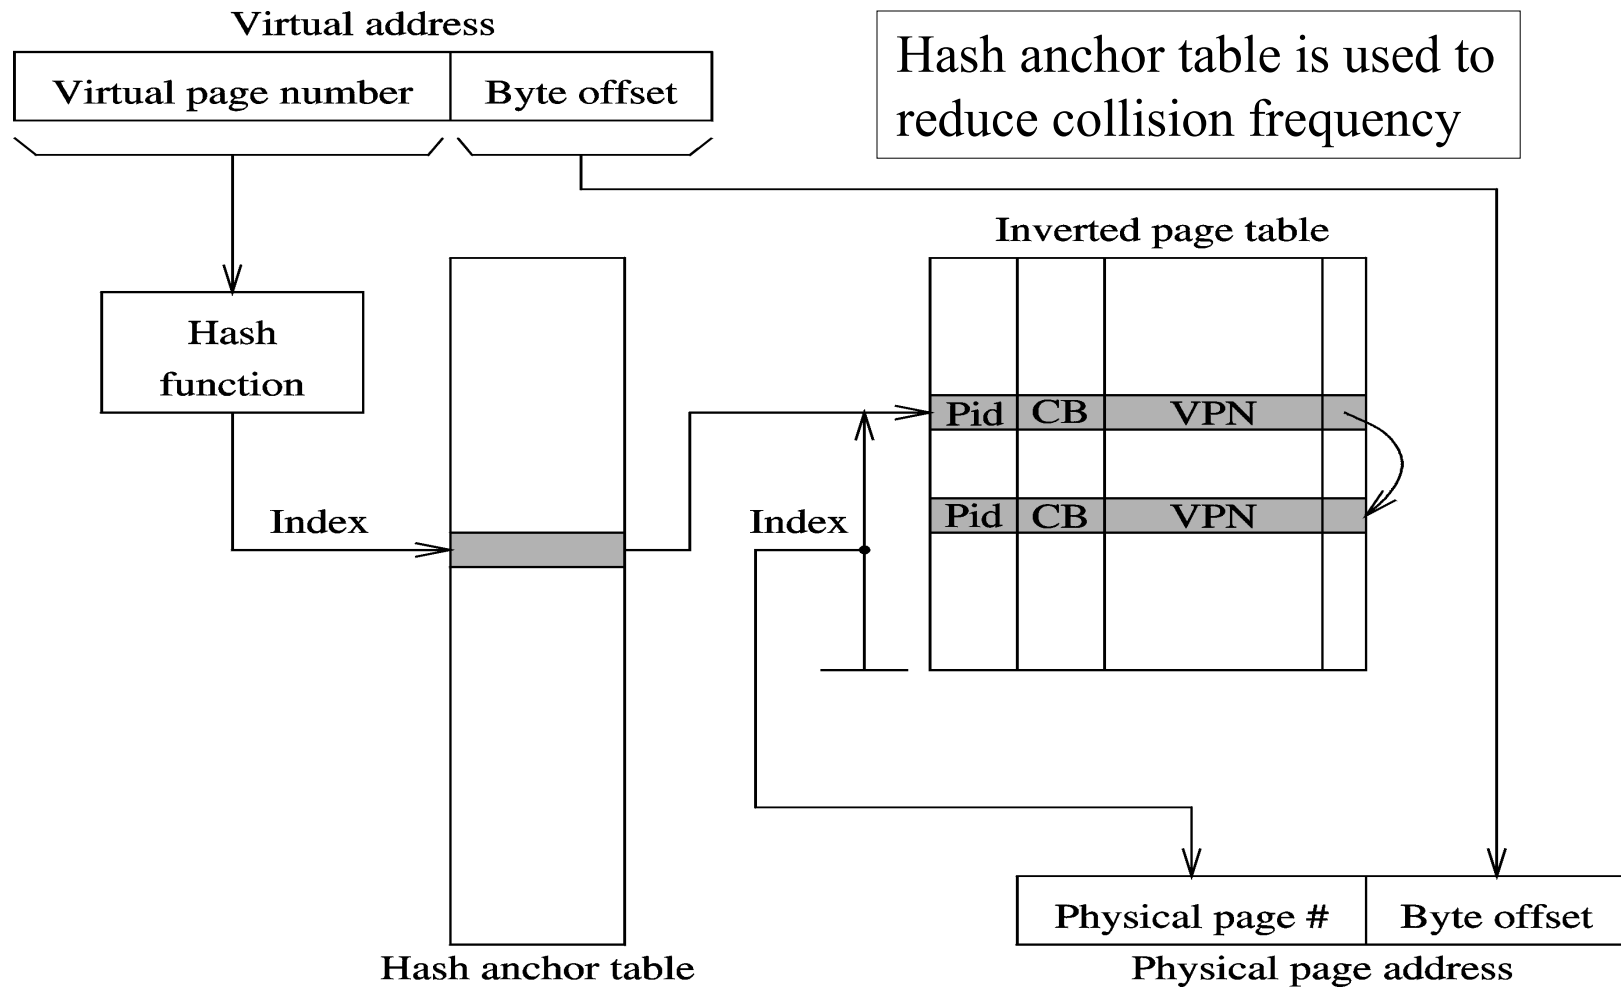
Page Table Placement (cont'd)

- Searching hierarchical page tables
 - * Two strategies
 - » Top-down
 - Starts at the root and follows all the levels
 - Previous example requires four memory accesses
 - ➔ 3 for the three page tables
 - ➔ 1 to read user data
 - » Bottom-up
 - Reduces the unacceptable overhead with top-down search
 - Starts at the bottom level
 - ➔ If the page is in memory, gets the physical page number
 - ➔ Else, resorts to top-down search

Inverted page Table Organization

- Number of entries in the forward page table
 - * Proportional to number of virtual pages
 - » Quite large for modern 64-bit processors
 - » Why?
 - We use virtual page number as index
- To reduce the number of entries
 - * Use physical page number as index
 - » Table grows with the size of memory
- Only one system-wide page table
 - * VPN is hashed to generate index into the table
 - » Needs to handle collisions

Inverted page Table Organization (cont'd)



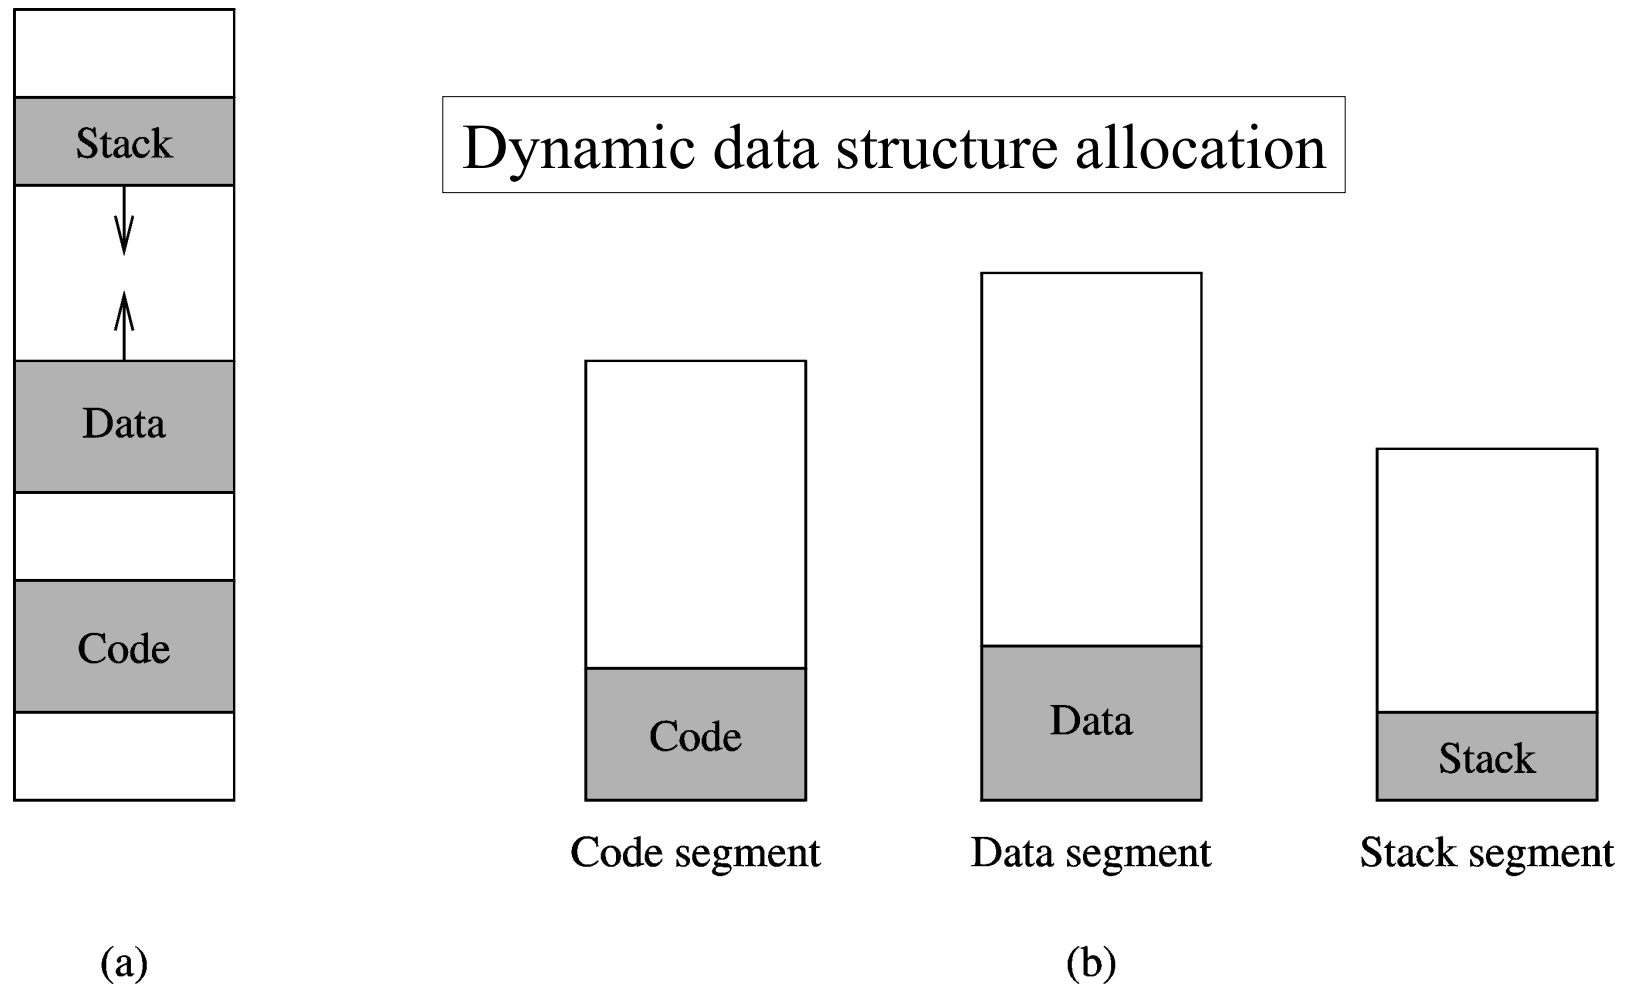
Segmentation

- Virtual address space is linear and 1-dimensional
 - * Segmentation adds a second dimension
- Each process can have several virtual address spaces
 - * These are called *segments*
 - * Example: Pentium
 - » Segments can be as large as 4 GB
- Address consists of two parts
 - * Segment number
 - * Offset within the segment

Segmentation (cont'd)

- Pentium and PowerPC support segmented-memory architecture
 - * Paging is transparent to programmer
 - » Segmentation is not
 - Pentium assembly programming makes it obvious
 - ➔ Uses three segments: data, code, and stack
- Segmentation offers some key advantages
 - * Protection
 - » Can be provided on segment-by-segment basis
 - * Multiple address spaces
 - * Sharing
 - » Segments can be shared among processes

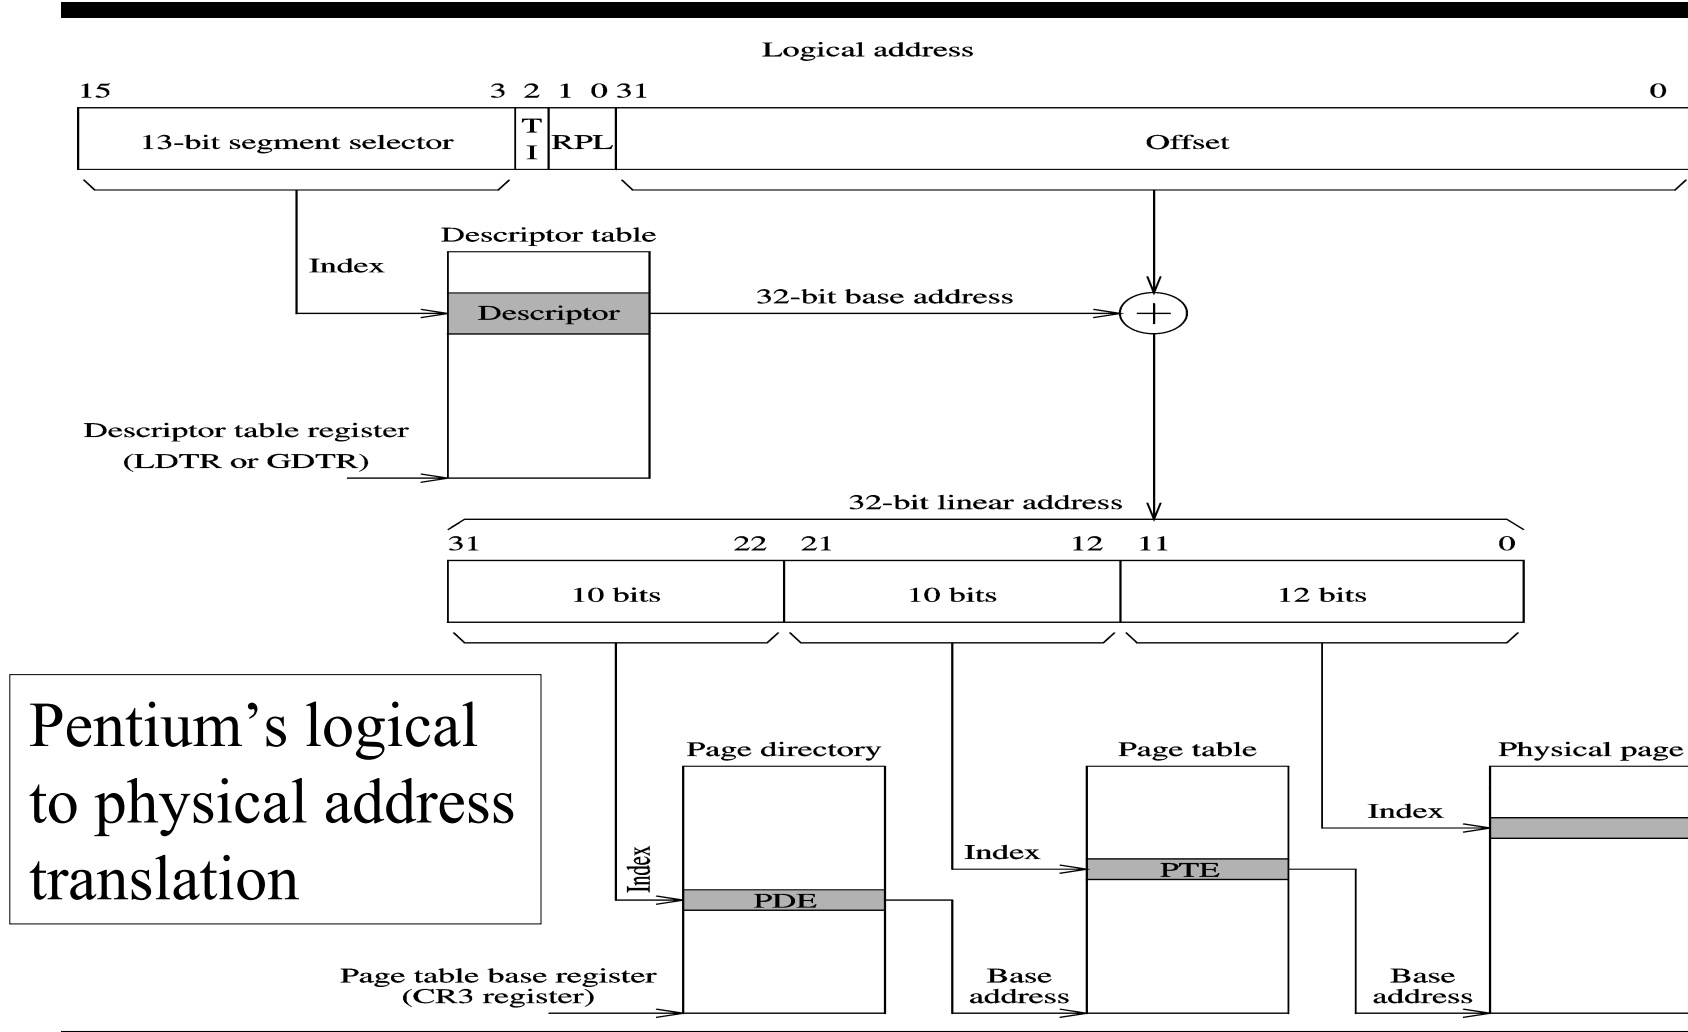
Segmentation (cont'd)



Example Implementations

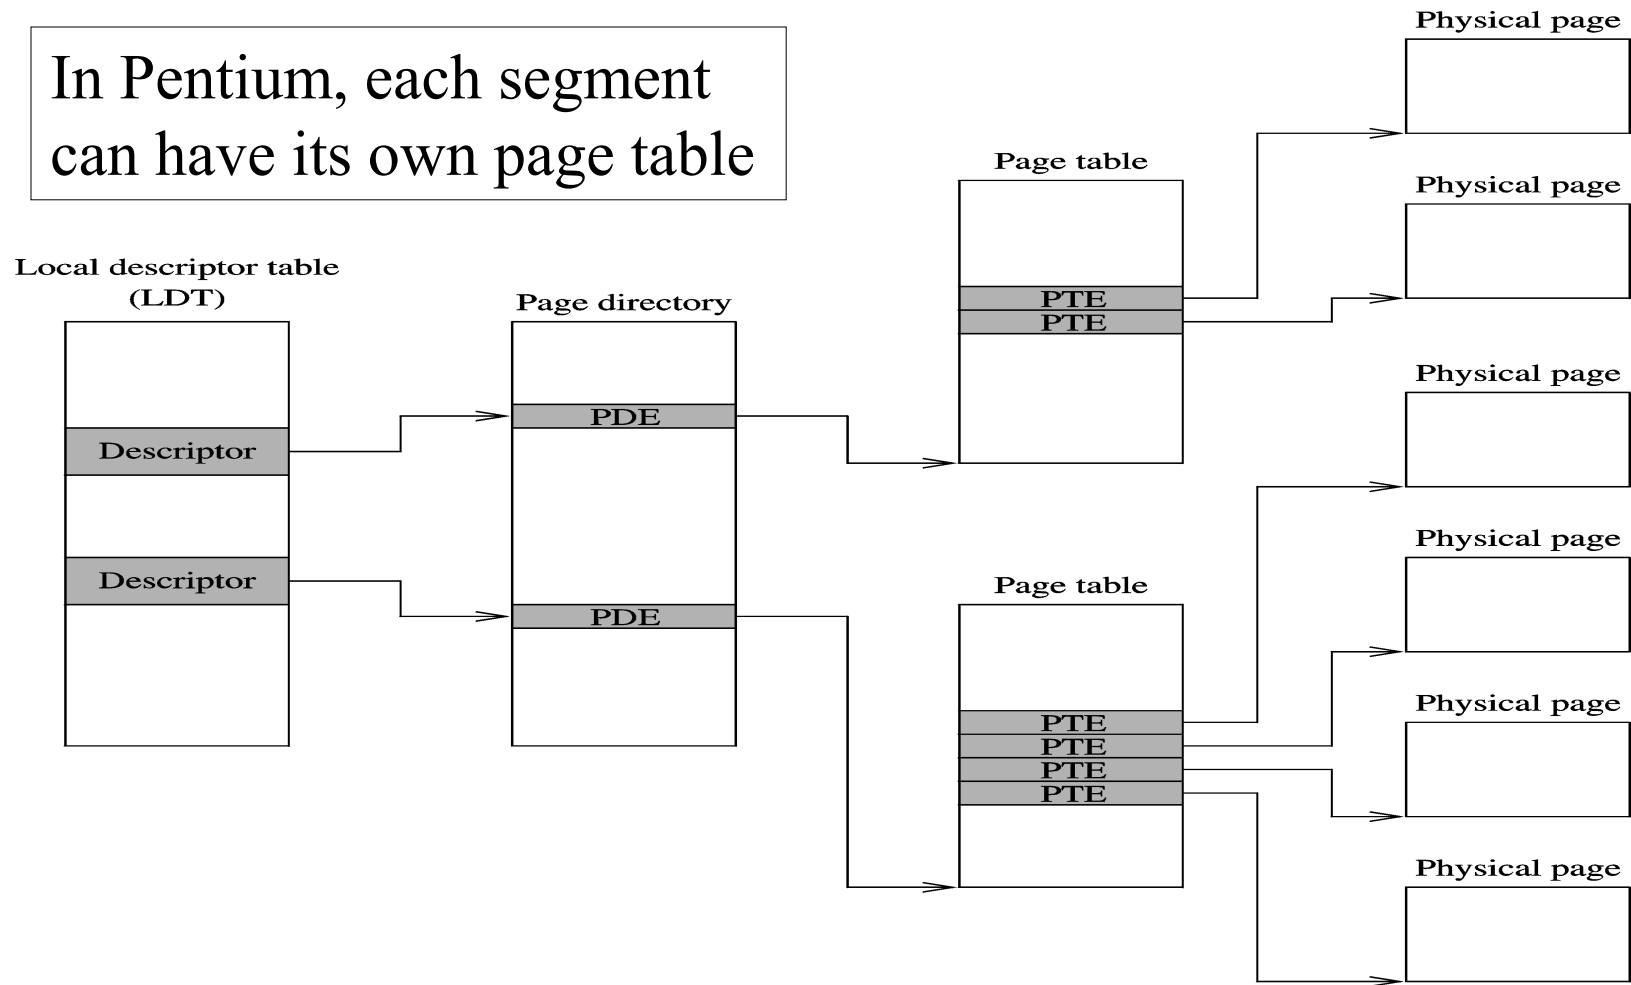
- Pentium
 - * Supports both paging and segmentation
 - » Paging can be turned off
 - » Segmentation can be turned off
 - * Segmentation translates a 48-bit logical address to 32-bit linear address
 - » If paging is used
 - It translates the 32-bit linear address to 32-bit physical address
 - » If paging is off
 - Linear address is treated as the physical address

Example Implementations (cont'd)



Example Implementations (cont'd)

In Pentium, each segment can have its own page table




Example Implementations (cont'd)



Pentium's page directory & page table entry format

(a) Page directory entry

0 for 4-KB page size



(b) Page table entry

Bit 0 Page present bit (P)

Bit 1 Writes permitted (W)

Bit 2 User/supervisor (U)

Bit 3 Page-level write-through (PWT)

Bit 4 Page-level cache-disable (PCD)

Bit 5 Page accessed (A)

Bit 6 0 in PDT
dirty bit (D) in PTE


Bit 7 Page size in PDE (0 for 4 KB pages)

Avail: Available for system programmer use

Shaded bits: Reserved by Intel (must be zero)

Example Implementations (cont'd)

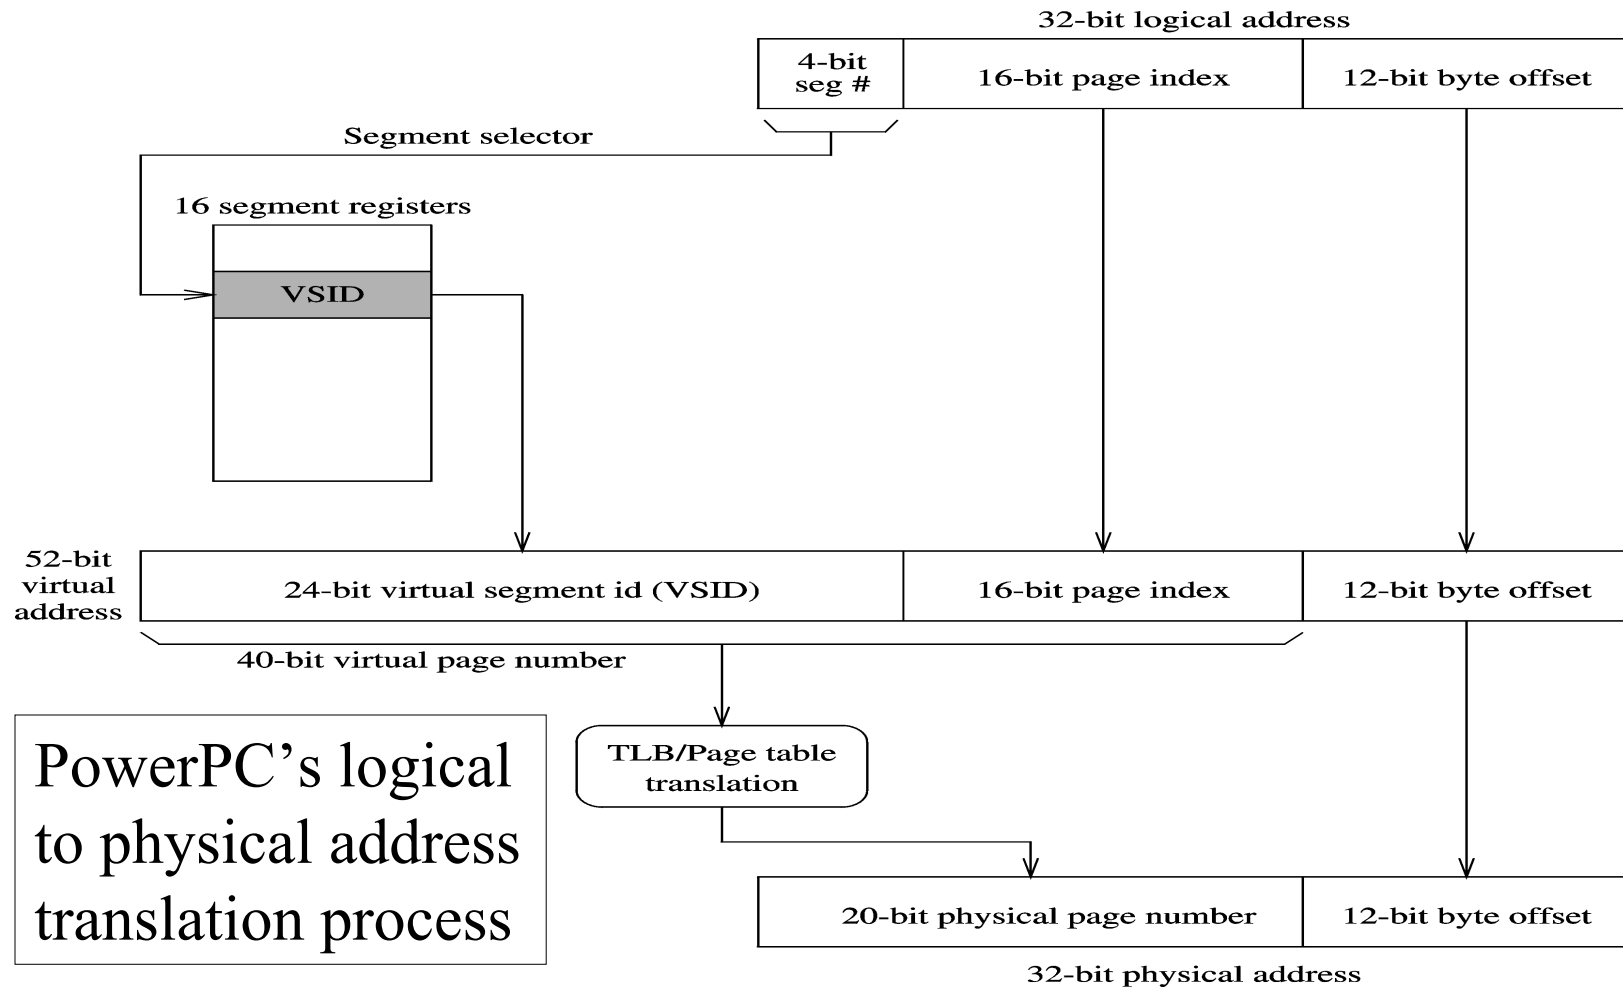
Pentium's
protection
rings



Example Implementations (cont'd)

- PowerPC
 - * Supports both segmentation and paging
 - * Logical and physical addresses are 32-bit long
 - * 32-bit logical address consists of
 - » 12-bit byte offset
 - » 16-bit page index
 - » 4-bit segment number
 - Selects one of 16 segment registers
 - Segment descriptor is a 24-bit virtual segment id (VSID)
 - * 52-bit virtual address consists of
 - » 40-bit VPN
 - » 12-bit offset

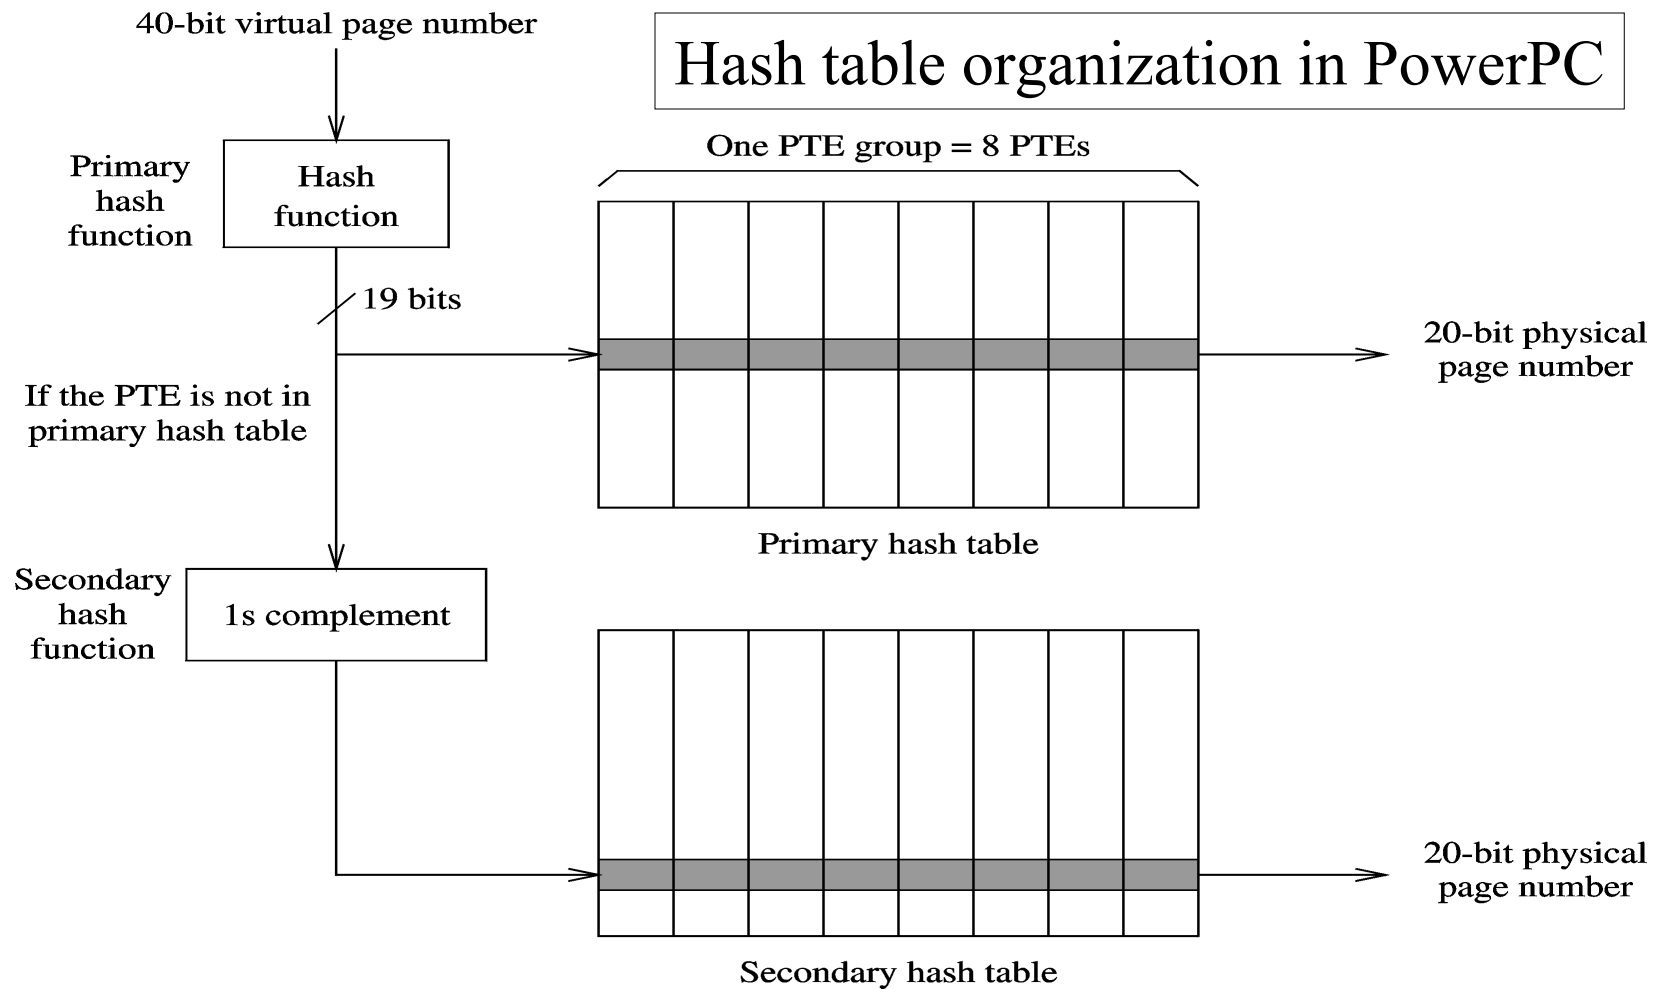
Example Implementations (cont'd)



Example Implementations (cont'd)

- PowerPC uses inverted page table
 - * Uses two hash tables
 - » Primary
 - Uses 8-way associative page table entry groups
 - » Secondary
 - 1s complement of the primary hash function
 - * PTEs are 8-bytes wide
 - » Stores valid bit, reference bit, changed bit (i.e., dirty bit)
 - » W bit (write-through)
 - $W = 1$: write-through policy
 - $W = 0$: write-back policy
 - » I bit (cache inhibit)
 - $I = 1$: cache inhibited (accesses main memory directly)

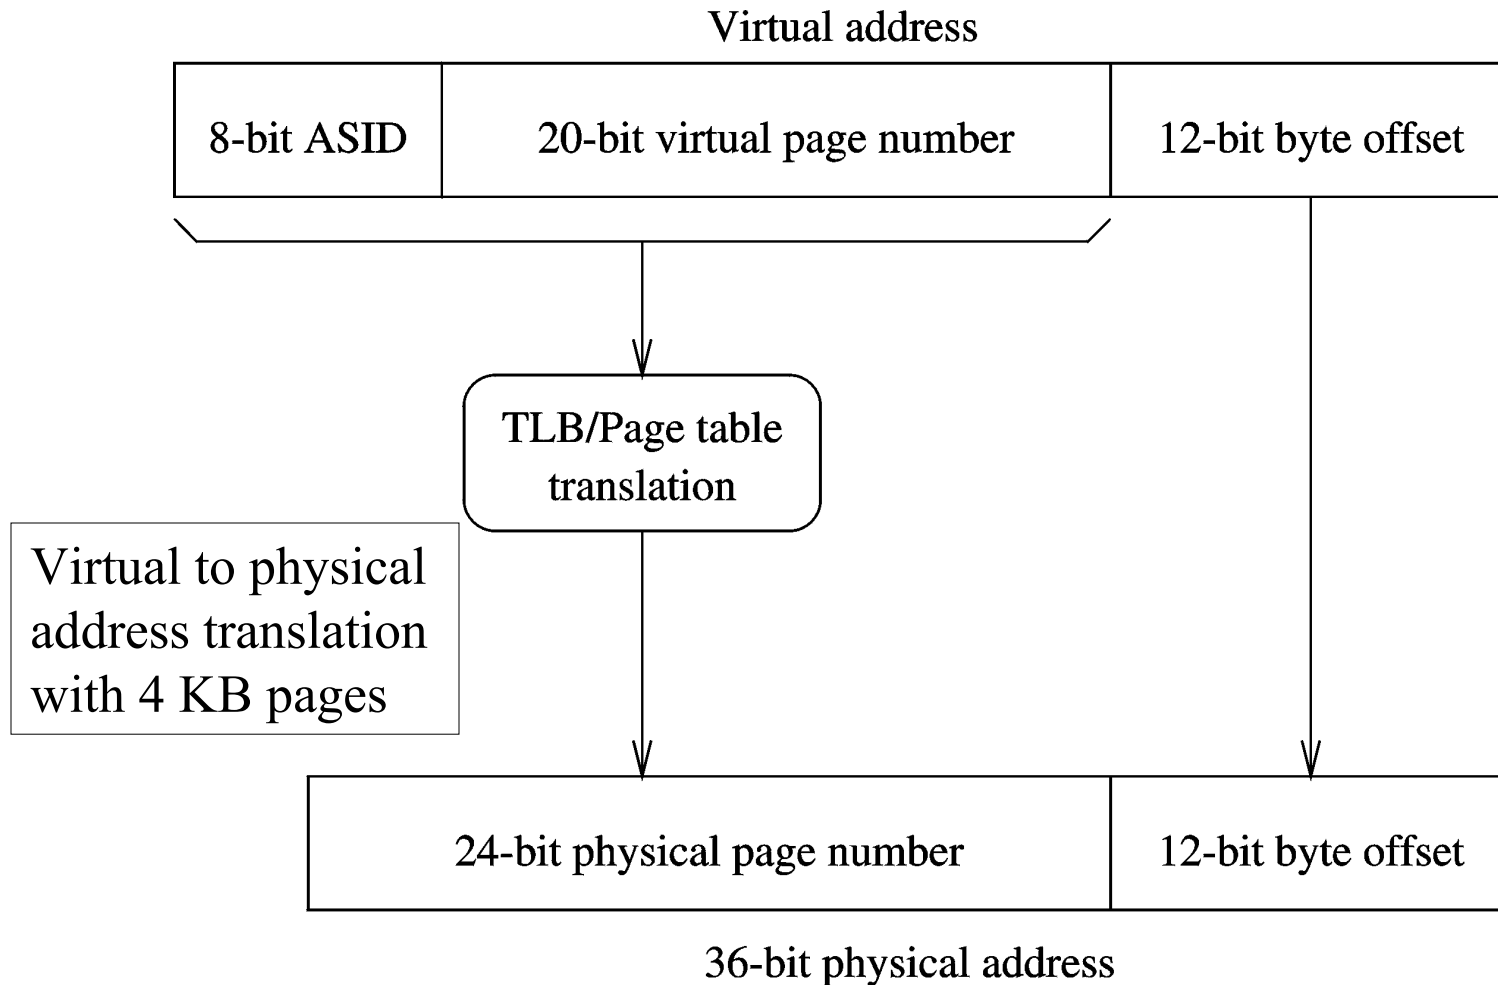
Example Implementations (cont'd)



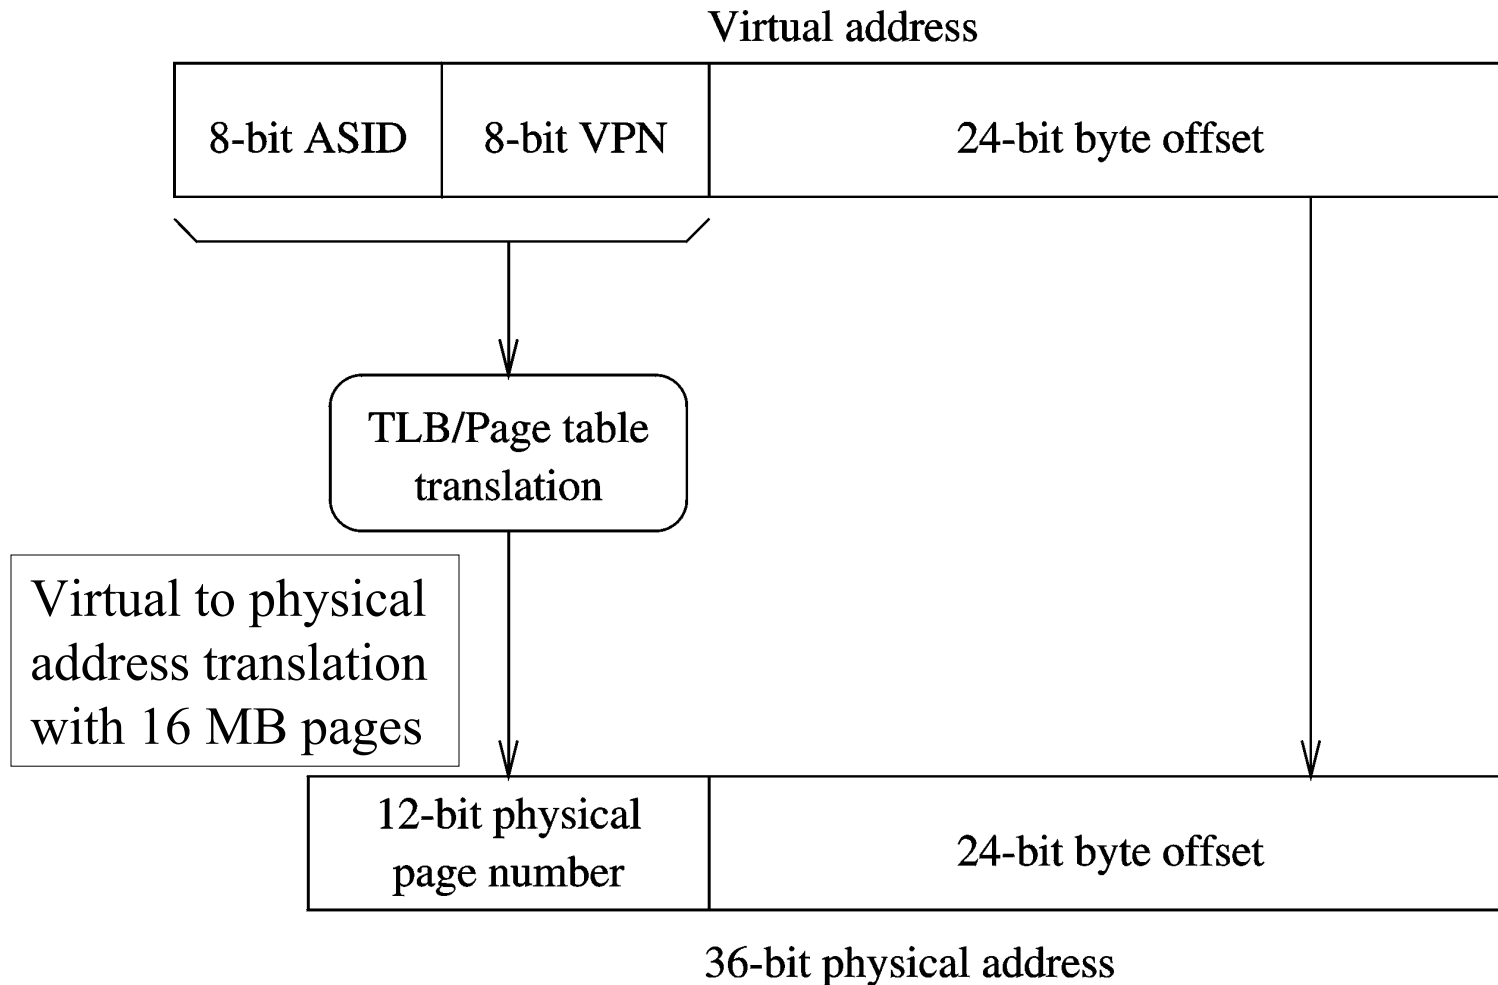
Example Implementations (cont'd)

- MIPS R4000
 - * Segmentation is not used
 - » Uses address space identifiers (ASIDs) for
 - Protection
 - Virtual address space extension
 - * 32-bit virtual address consists of
 - » 8-bit ASID
 - » 20-bit VPN
 - Depends on the page size
 - » 12-bit offset
 - Depends on the page size
 - Supports pages from 4 KB to 16 MB

Example Implementations (cont'd)




Example Implementations (cont'd)

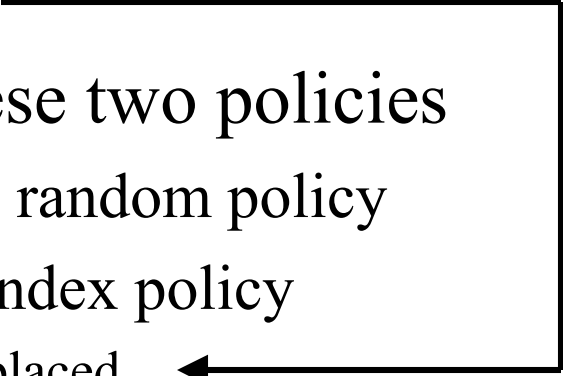


Example Implementations (cont'd)

MIPS TLB entry format



Example Implementations (cont'd)

- MIPS R 4000 supports two TLB replacement policies
 - * Random
 - » Randomly selects an entry
 - * Indexed
 - » Selects the entry specified
 - Two registers support these two policies
 - * A Random register for the random policy
 - * An Index register for the index policy
 - » Specifies the entry to be replaced
- 

Last slide